



Can you create some really good information about this project to keep it in my portfolio

AgentOS is an **MCP-native operating system** for autonomous multi-agent LLM workflows, designed as a research-grade platform that makes your agents look and feel like OS processes running over a standardized tool and memory stack. Below is portfolio-ready content you can paste into your website, CV, LinkedIn, or GitHub.[\[1\]](#) [\[2\]](#)

One-line project pitch

AgentOS – An MCP-native operating system that treats LLM agents as processes and MCP tools as system resources to enable secure, scalable, and long-running autonomous multi-agent workflows.[\[2\]](#) [\[1\]](#)

Project overview

AgentOS is a research-focused operating system for AI agents, built entirely on the Model Context Protocol (MCP) to provide a unified environment for tool access, memory, scheduling, and safety. Instead of wiring together ad-hoc tools and agents, AgentOS introduces OS-style abstractions such as a kernel, drivers, processes, and hierarchical memory for LLM-based systems.[\[1\]](#) [\[2\]](#)

The system targets core pain points in current agent frameworks: fragmented context, non-standard tool APIs, unreliable long-running behavior, and the lack of a central “agent kernel” to manage processes, resources, and safety policies. By framing the entire multi-agent stack as an operating system, AgentOS is designed to be both a practical engineering platform and a strong foundation for academic publication in AI systems and autonomous agents.[\[2\]](#) [\[1\]](#)

Key technical contributions

- **OS-style agent kernel**
 - Central coordinator that manages agent lifecycles, schedules tasks, routes context, and enforces tool permissions, analogous to an operating system kernel for LLM agents.[\[1\]](#) [\[2\]](#)
 - Maintains system-wide logs and metrics to support observability, debugging, and research analysis.[\[2\]](#) [\[1\]](#)
- **MCP tool driver layer**

- Treats MCP tools as system drivers (filesystem, browser/HTTP client, vector memory store, database client, logging/audit, evaluation harness) exposed through a standardized interface.[\[1\]](#) [\[2\]](#)
- Encourages modularity so new tools can be plugged in without changing core agent orchestration logic.[\[2\]](#) [\[1\]](#)

- **Multi-agent process layer**

- Models agents as specialized processes (retrieval, planning, reasoning, execution, safety/review) that cooperate through the kernel rather than as isolated, hard-coded pipelines.[\[1\]](#) [\[2\]](#)
- Enables cooperative multitasking, clearer separation of responsibilities, and safer inter-agent communication.[\[2\]](#) [\[1\]](#)

- **Hierarchical memory system**

- Implements multi-tier memory inspired by OS caches:
 - L1: short-term dialogue buffer for immediate context.[\[1\]](#) [\[2\]](#)
 - L2: episodic vector memory for semantic retrieval.[\[2\]](#) [\[1\]](#)
 - L3: structured JSON schemas for workflows and state.[\[1\]](#) [\[2\]](#)
 - L4: cold storage via MCP filesystem for long-term archival.[\[2\]](#) [\[1\]](#)
- Designed to maintain state across long-running tasks while controlling context size and retrieval accuracy.[\[1\]](#) [\[2\]](#)

- **Zero-trust tool firewall and observability**

- Enforces a zero-trust policy where every tool call is validated, risk-scored, and logged for traceability.[\[2\]](#) [\[1\]](#)
- Tracks metrics such as latency, token usage, memory hits/misses, and safety events to support rigorous evaluation.[\[1\]](#) [\[2\]](#)

Architecture at a glance

- **Agent Kernel (Coordinator)**

- Schedules agent processes, routes messages and context, enforces permissions, and keeps global system logs.[\[2\]](#) [\[1\]](#)

- **MCP Tool Layer (Drivers)**

- Filesystem driver for persistent data, browser/HTTP for web access, vector store for embeddings, database client for structured data, and logging/audit tools.[\[1\]](#) [\[2\]](#)

- **Multi-Agent Layer**

- Retrieval Agent: interfaces with tools and memory to gather relevant information.[\[2\]](#) [\[1\]](#)
- Planning Agent: decomposes high-level tasks and coordinates downstream agents.[\[1\]](#) [\[2\]](#)
- Reasoning Agent: performs inference, checking plans and outputs.[\[2\]](#) [\[1\]](#)

- Execution Agent: triggers tools and carries out actions in the environment.[\[1\]](#) [\[2\]](#)
- Safety/Review Agent: audits actions and outputs against policy and risk criteria.[\[2\]](#) [\[1\]](#)

- **Memory and Safety**

- Multi-layer memory architecture ensures scalable context management with controlled access and retrieval.[\[1\]](#) [\[2\]](#)
- Zero-trust firewall validates each tool interaction, with logging and auditing for reproducible experiments.[\[2\]](#) [\[1\]](#)

Research focus and evaluation

- **Research objectives**

- Design an OS-style architecture for agents using MCP and multi-agent orchestration methods.[\[3\]](#) [\[2\]](#)
- Build a modular MCP tool ecosystem, implement an Agent Scheduler, and construct hierarchical memory and a trust/safety firewall.[\[3\]](#) [\[2\]](#)

- **Evaluation metrics**

- Performance: tool-call latency, inter-agent communication overhead, memory retrieval accuracy, task completion time.[\[1\]](#) [\[2\]](#)
- Reliability: state consistency during long-running workflows, memory tier hit rates, and recovery from agent failures.[\[2\]](#) [\[1\]](#)
- Safety: blocked unsafe tool calls, policy violation detection accuracy, and false positive/negative rates in auditing.[\[1\]](#) [\[2\]](#)

- **Comparative analysis**

- Benchmarks AgentOS against LangChain Agents, LangGraph workflows, and CrewAI multi-agent systems.[\[2\]](#) [\[1\]](#)
- Highlights novelty in OS-style coordination, hierarchical memory, zero-trust tool governance, and detailed observability beyond typical workflow-based frameworks.[\[1\]](#) [\[2\]](#)

What this project demonstrates about you

- Ability to architect and implement a complex OS-inspired multi-agent system on top of MCP, rather than building a simple single-agent app.[\[2\]](#) [\[1\]](#)
- Hands-on experience with protocol design, multi-agent orchestration, hierarchical memory, and safety-aware tool usage, with a clear evaluation and benchmarking plan suitable for a workshop or conference paper.[\[1\]](#) [\[2\]](#)

If you want, a next step can be writing a short “Project Highlights” section tailored specifically for LinkedIn or a portfolio website card (2–3 sentences plus 3 bullet points).

With what can you compare this project then? How is this going to work on windows? or linux? or is it a software app or it is similiar to ubuntu ?

AgentOS is conceptually similar to an operating system **for AI agents**, but in practice it runs as a software service on top of Windows, Linux, or the cloud, not as a bare-metal OS like Ubuntu.^[4]
^[5]

What you can compare it to

For your portfolio and explanations, you can compare AgentOS to:

- LangChain Agents and LangGraph
 - These provide agent workflows and tool calling, but they do not define a kernel, drivers, or hierarchical memory as OS-style abstractions.^[5] ^[4]
- CrewAI and other multi-agent orchestrators
 - They coordinate multiple agents, but usually with flatter context and simpler security models, not a structured kernel + zero-trust firewall design.^[4] ^[5]
- An “application-layer OS”
 - AgentOS borrows ideas from operating systems (processes, drivers, kernel, memory tiers) and applies them to LLM agents using MCP.^[5] ^[4]

Here's a simple comparison table you can use:

Aspect	AgentOS (your project)	LangChain / LangGraph / CrewAI	Ubuntu / Windows (real OS)
Main purpose	OS-style coordination of LLM agents and tools ^[4] ^[5]	Agent workflows and tool integrations ^[4] ^[5]	Manage hardware, processes, user apps
Runs on	On top of an existing OS (Windows, Linux, cloud) ^[4] ^[5]	On top of an existing OS	Directly on hardware
“Kernel”	Agent kernel for scheduling and permissions ^[4] ^[5]	No true kernel, mostly orchestrator code	Full OS kernel managing CPU, memory, devices
“Processes”	Agents as processes (retrieval, planning, etc.) ^[4] ^[5]	Agents/tasks but not treated like OS processes	Native processes and threads
Memory model	Multi-tier (L1-L4) hierarchical memory ^[4] ^[5]	Mostly flat context, basic memories ^[4]	RAM, caches, disk, virtual memory
Tool access	MCP tools as drivers + zero-trust firewall ^[4] ^[5]	Tool calls with simpler access controls ^[4]	Syscalls, drivers, and OS permissions

Is it like Ubuntu or Windows?

- Not a full OS like Ubuntu/Windows
 - AgentOS does not replace your machine's operating system or boot on bare metal.^[4]
^[5]
 - It is "OS-style" at the application layer for agents: the kernel is a coordinator service, and the drivers are MCP tools, not hardware device drivers.^[5] ^[4]
- What "OS-style" means here
 - You have a central **agent kernel** managing processes (agents), permissions, scheduling, and logs, just like an OS kernel manages programs.^[4] ^[5]
 - You have a **hierarchical memory system** analogous to L1/L2 caches vs disk, but implemented as dialogue buffer, vector store, JSON state, and cold storage.^[5] ^[4]

How it works on Windows or Linux in practice

You can describe it as a cross-platform software system:

- Deployment model
 - AgentOS is implemented as a software application (or service) that runs on top of any OS that can host your MCP server and associated services (Python backend, vector DB, etc.).^[4] ^[5]
 - On Windows or Linux, you would typically run it as:
 - A long-running backend process (e.g., Python/Node service)
 - Connected MCP tools (filesystem, browser client, vector store, DB)
 - An LLM endpoint (local or cloud) wired into the kernel.^[5] ^[4]
- Interaction model
 - Users or higher-level applications send tasks to AgentOS.
 - The agent kernel spins up or routes to agent "processes" (retrieval, planning, execution, safety).^[4] ^[5]
 - Those agents call MCP tools (drivers) under the rules of the zero-trust firewall and use the hierarchical memory for context.^[5] ^[4]

So, for your portfolio you can say something like:

"AgentOS is not a replacement for Ubuntu or Windows. It is an OS-style runtime that runs on top of any standard OS and provides kernel-like coordination, drivers, and hierarchical memory specifically for LLM agents using MCP."^[4] ^[5]

If you want, I can draft a short "FAQ" section (3–5 Q&A) you can add to your portfolio page to answer exactly these questions for reviewers.

Give me a perfect example usecase, where does this project work in real life. How much useful is this idea

A strong real-world use case is a 24/7 “AI operations control room” where multiple agents monitor systems, react to incidents, and coordinate actions across tools more reliably than a single prompt-engineered bot.^[7] ^[8]

Concrete example use case

Imagine a large SaaS company running an always-on AI operations assistant:

- A **Monitoring Agent** streams logs, metrics, and alerts from monitoring tools through MCP drivers (filesystem, HTTP, DB), and detects anomalies.^[8] ^[7]
- A **Retrieval Agent** pulls related past incidents, runbooks, tickets, and code changes from vector memory and structured JSON memory.^[7] ^[8]
- A **Planning Agent** creates a step-by-step response plan (who to notify, which checks to run, what to roll back), coordinating with other agents via the kernel.^[8] ^[7]
- An **Execution Agent** safely triggers scripts, API calls, and rollbacks through tool drivers, only for actions that pass the zero-trust firewall policies.^[7] ^[8]
- A **Safety/Review Agent** audits each action, logs decisions, and blocks risky tool calls, ensuring everything is traceable.^[8] ^[7]

Because AgentOS has hierarchical memory, it can remember previous incidents (L2/L3/L4) and improve over time, while still keeping the active context small and efficient. This is hard to do with a single agent + flat context, but natural when you have an OS-style kernel, processes, and memory tiers.^[7] ^[8]

You can pitch it like this in your portfolio:

“In production, AgentOS could power a 24/7 AI Ops assistant where specialized agents monitor systems, retrieve historical context, plan remediation, and execute safe actions under a zero-trust firewall, with all activity logged and auditable.”

Where this idea is useful

The proposal explicitly calls out several domains where long-running, tool-using, multi-agent systems with strong memory and safety are important:^[8] ^[7]

- **Robotics** – coordinating perception, planning, and control agents that share memory and tools safely.^[8]
- **Autonomous reasoning systems** – complex tasks broken into retrieval, planning, reasoning, and execution agents.^[7] ^[8]

- **Real-time multi-agent coordination** – for logistics, fleet management, or incident response. [8]
- **Industrial automation** – where actions must be auditable, gated by policy, and robust over long periods. [7] [8]
- **AI personal assistants** – that persist long-term memory and use many tools (calendar, email, files, web, devices) with fine-grained safety controls. [8]

How useful is the idea, realistically?

From the proposal itself, the idea is framed as both practically and academically valuable: [9] [7]
[8]

- It solves real limitations: fragmented context, proprietary tool APIs, brittle long-running agents, and the lack of an “agent kernel.” [8]
- It introduces **new structure**: OS-style kernel, drivers, hierarchical memory, and zero-trust tool governance for agents, which mainstream frameworks do not provide today. [9] [7] [8]
- It is designed for **autonomous, long-running workflows**, which is where current agents often fail due to context overflow, state inconsistency, and weak safety. [8]
- It is positioned as a **publishable contribution** suitable for workshops or conferences in AI systems, LLM engineering, and autonomous agents. [9] [7] [8]

A portfolio-friendly way to state the usefulness:

“AgentOS targets the next generation of AI systems: long-running, tool-using, multi-agent workflows that must be reliable, stateful, and safe. It provides structure (kernel, drivers, memory tiers, firewall) that today’s agent frameworks are missing, making it a practically useful and academically novel idea.” [7] [8]

If you’d like, I can help you turn this into a 3–4 line “Real-world impact” block you can drop directly into your portfolio page.

**

Can you make all of this chat in an MD file

Below is a Markdown version you can copy into an .md file. It’s written as a clean portfolio page, not as a chat log.

AgentOS: An MCP-Native Operating System for Autonomous Multi-Agent AI Systems

AgentOS is an **MCP-native operating system** for autonomous multi-agent LLM workflows. It treats LLM agents as processes and MCP tools as system resources, providing an OS-style environment for secure, scalable, long-running autonomous AI systems.

1. One-Line Project Pitch

AgentOS – An MCP-native operating system that treats LLM agents as processes and MCP tools as system resources to enable secure, scalable, and long-running autonomous multi-agent workflows.

2. Project Overview

Large Language Models (LLMs) are evolving from passive chatbots into active, tool-using agents that perform retrieval, reasoning, planning, and real-world actions. Current multi-agent frameworks often wire these pieces together in an ad-hoc way, without a unified “operating system” abstraction.

AgentOS proposes an OS-style architecture for AI agents built natively on the Model Context Protocol (MCP). Instead of just building a workflow or a single agent, it introduces concepts such as a kernel, processes, drivers, and hierarchical memory—specifically for LLM-based systems.

The project targets four core limitations of today’s agent stacks:

- Fragmented context and siloed tools.
- Non-standard, proprietary tool APIs.
- Unreliable long-running agents (context overflow, state loss).
- No unified “agent kernel” to handle scheduling, memory, permissions, and safety.

By framing the multi-agent system as an operating system, AgentOS aims to be both a practical engineering platform and a strong foundation for research in AI systems and autonomous agents.

3. Key Technical Contributions

3.1 Agent Kernel (OS-Style Coordinator)

- Central coordinator that manages agent lifecycles and scheduling.
- Routes context and messages between agents.
- Enforces tool permissions and security policies.
- Maintains global logs and metrics for observability and research.

This kernel plays a role similar to an OS kernel, but for LLM agents instead of native processes.

3.2 MCP Tool Driver Layer

- Treats MCP tools as “drivers” for system resources:
 - Filesystem (persistent storage).
 - Browser/HTTP client (web access).
 - Vector memory store.

- Database client.
- Logging and audit tools.
- Evaluation harness.
- Standardized MCP interfaces make tools modular and swappable.
- Core orchestration logic does not depend on any single tool implementation.

3.3 Multi-Agent Process Layer

- Models agents as processes with clear responsibilities:
 - Retrieval Agent.
 - Planning Agent.
 - Reasoning Agent.
 - Execution Agent.
 - Safety/Review Agent.
- Agents run cooperatively via the kernel, not as isolated, hard-coded chains.
- Enables cleaner separation of concerns and safer inter-agent communication.

3.4 Hierarchical Memory System

AgentOS implements multi-tier memory inspired by OS memory hierarchies:

- **L1 – Dialogue Buffer**
Short-term conversation memory for immediate context.
- **L2 – Episodic Memory (Vector DB)**
Embedding-based episodic memory for semantic retrieval.
- **L3 – Structured Memory (JSON)**
Structured schemas to track workflows, state machines, and important facts.
- **L4 – Cold Storage (Filesystem)**
Long-term, infrequently accessed data stored via MCP filesystem tools.

This design allows long-running workflows to maintain rich state without overflowing a single context window.

3.5 Zero-Trust Firewall and Observability

- Zero-trust firewall validates every tool call.
- Risk scoring and policy checks gate access to tools and data.
- Comprehensive logging for reproducibility and auditing.
- Metrics collected include:
 - Tool-call latency.
 - Inter-agent communication overhead.

- Memory hit/miss rates.
- Task completion time.
- Safety events (blocked calls, violations).

4. Architecture Overview

4.1 High-Level Components

- **Agent Kernel (Coordinator)**
Manages agent lifecycles, scheduling, context routing, permissions, and system logs.
- **MCP Tool Drivers (Resource Layer)**
Filesystem, HTTP client, vector store, database, logging/audit, and evaluation harness exposed via MCP.
- **Multi-Agent Process Layer**
Specialized agents (retrieval, planning, reasoning, execution, safety) that cooperate through the kernel.
- **Hierarchical Memory System**
L1–L4 memory tiers mapping to short-term context, episodic embeddings, structured state, and cold storage.
- **Zero-Trust Firewall**
Policy-driven gating of all tool calls with risk scoring and logging.
- **Monitoring & Observability**
Detailed diagnostics for performance, reliability, and safety.

5. Comparison to Other Systems

5.1 Conceptual Comparison

AgentOS can be viewed as an “application-layer operating system” for LLM agents:

- It does **not** replace Ubuntu or Windows.
- It **runs on top of** an existing OS (Windows, Linux, cloud instance).
- It introduces OS-like abstractions (kernel, drivers, processes, memory tiers) at the agent level.

5.2 Comparison Table

Aspect	AgentOS (this project)	LangChain / LangGraph / CrewAI	Ubuntu / Windows (real OS)
Main purpose	OS-style coordination of LLM agents and tools	Agent workflows and tool integrations	Manage hardware, processes, and user apps
Runs on	On top of an existing OS	On top of an existing OS	Directly on hardware

Aspect	AgentOS (this project)	LangChain / LangGraph / CrewAI	Ubuntu / Windows (real OS)
"Kernel"	Agent kernel for scheduling and permissions	Orchestrator code, no OS-style kernel	Full OS kernel managing CPU, memory, devices
"Processes"	Agents as processes (retrieval, planning, etc.)	Agents/tasks, not treated like OS processes	Native processes and threads
Memory model	Multi-tier hierarchical memory (L1-L4)	Mostly flat context/memory	RAM, caches, disk, virtual memory
Tool access	MCP tools as drivers + zero-trust firewall	Tool calls with simpler access controls	Syscalls, drivers, and OS permissions

6. Real-World Use Case Example

6.1 24/7 AI Operations Control Room (AI Ops)

A strong real-world scenario for AgentOS is a 24/7 AI operations assistant for a SaaS or cloud company.

Agents and responsibilities:

- **Monitoring Agent**

Continuously reads logs, metrics, and alerts via MCP tools (filesystem, HTTP, DB). Detects anomalies and incidents.

- **Retrieval Agent**

Uses hierarchical memory to pull:

- Past incidents.
- Runbooks and documentation.
- Historical tickets and relevant code changes.

- **Planning Agent**

Creates a remediation plan:

- Diagnostics to run.
- Services to restart or roll back.
- Stakeholders to notify.

- **Execution Agent**

Safely triggers scripts and API calls:

- Only for actions that pass firewall rules.
- Uses MCP tools to interact with infrastructure.

- **Safety/Review Agent**

Audits each action:

- Checks against policies.

- Blocks risky actions.
- Ensures all steps are logged and reproducible.

Why AgentOS is a good fit here:

- Long-running workflows need persistent, structured memory across many incidents.
- Multiple specialized agents must coordinate safely over shared tools.
- Every action must be auditable and policy-compliant.
- OS-style abstractions (kernel, processes, drivers, memory tiers) reduce complexity and make behavior easier to reason about.

Portfolio-ready description:

In production, AgentOS could power a 24/7 AI Ops assistant where specialized agents monitor systems, retrieve historical context, plan remediation, and execute safe actions under a zero-trust firewall, with all activity logged and auditable.

7. Where AgentOS is Useful

AgentOS is particularly relevant for domains that require long-running, tool-using, multi-agent systems with strong safety and memory:

- Robotics and autonomous systems.
- Real-time multi-agent coordination (logistics, fleet management).
- Industrial automation with strict safety and compliance needs.
- Complex autonomous reasoning systems (research assistants, workflow managers).
- Advanced AI personal assistants that must remember user context for months, interact with many tools, and remain safe.

8. Practical Deployment Model

AgentOS is implemented as a **software runtime**, not a bare-metal OS:

- Runs on top of Windows, Linux, or cloud environments.
- Consists of:
 - An agent kernel service (e.g., Python/Node backend).
 - MCP tools (filesystem, HTTP client, vector DB, relational DB, logging).
 - LLM endpoints (local or cloud).
- Clients (users, apps, services) send tasks to AgentOS; the kernel coordinates agents and tools to complete them.

You can summarize it like this:

AgentOS is not a replacement for Ubuntu or Windows. It is an OS-style runtime that runs on top of any standard OS and provides kernel-like coordination, drivers, hierarchical memory, and a zero-trust firewall specifically for LLM agents.

9. Research Focus and Evaluation

9.1 Research Objectives

- Design an OS-style architecture for multi-agent LLM systems using MCP.
- Build a modular MCP tool ecosystem (filesystem, search, web, DB, logging, vector memory).
- Implement an agent scheduler for cooperative multi-agent processes.
- Construct hierarchical memory layers (short-term, episodic, structured, cold storage).
- Add a risk-scored trust and safety firewall for tool usage.
- Benchmark performance, reliability, and autonomy against existing frameworks.

9.2 Evaluation Metrics

- **Performance**
 - Tool-call latency.
 - Inter-agent communication overhead.
 - Memory retrieval accuracy.
 - Task completion time.
- **Reliability**
 - State consistency in long-running workflows.
 - Memory tier hit rates.
 - Recovery after agent failure.
- **Safety**
 - Number of blocked unsafe tool calls.
 - Policy violation detection accuracy.
 - False positive/negative rates for auditing.
- **Comparative Analysis**
 - Benchmarked against LangChain Agents, LangGraph workflows, and CrewAI multi-agent systems.

10. Impact and Usefulness

AgentOS is useful at two levels:

1. Practical impact

- Addresses real problems in current agent frameworks:
 - Fragmented context and tool silos.
 - Proprietary, non-standard tool interfaces.
 - Fragility of long-running agents.
- Provides structure (kernel, drivers, memory tiers, firewall) that makes it easier to build robust, auditable, production-grade multi-agent systems.

2. Academic and career impact

- Bridges OS concepts (processes, drivers, kernel, memory hierarchy) with LLM agent systems—an underexplored research area.
- Suitable for workshop or conference papers in AI systems, LLM engineering, and autonomous agents.
- Demonstrates deep understanding of:
 - Multi-agent orchestration.
 - Protocol design (MCP).
 - Hierarchical memory.
 - Safety and observability in tool-using agents.

Portfolio-style positioning:

AgentOS targets the next generation of AI systems: long-running, tool-using, multi-agent workflows that must be reliable, stateful, and safe. By introducing OS-style abstractions—kernel, drivers, hierarchical memory, and a zero-trust firewall—it fills a gap that today's agent frameworks do not fully address, making it both practically useful and academically novel.

11. What This Project Shows About Me

- Ability to architect and implement complex, OS-inspired multi-agent systems on top of modern AI protocols.
- Hands-on experience with:
 - MCP integration.
 - Multi-agent orchestration.
 - Hierarchical memory design.
 - Safety-aware tool governance.
 - Benchmarking and evaluation.

- Readiness to explain and defend system design decisions in an academic or industry setting (e.g., for research roles or applied AI engineering roles).

If you want, I can next:

- Trim this down into a shorter "Project card" version for your portfolio homepage, and
- A very concise 2–3 line version for your resume or LinkedIn "Experience / Projects" section.

**

1. AgentOS.pdf
2. Project-Proposal-Report.pdf
3. AgentOS_Proposal.pdf
4. AgentOS.pdf
5. Project-Proposal-Report.pdf
6. AgentOS_Proposal.pdf
7. AgentOS.pdf
8. Project-Proposal-Report.pdf
9. AgentOS_Proposal.pdf
10. AgentOS.pdf
11. AgentOS_Proposal.pdf
12. Project-Proposal-Report.pdf